

## Description of Lab 1:

Lab 1 was implementing SimpleDB, a relational database that stores tables with tuples as each row. These were represented with the `Tuple` class which stores each `Field` as a list whose types are determined by the `TupleDesc`. It stores the metadata for each `Tuple` including the names and types of each `Field`.

At the heart of SimpleDB is the `Database` class that is a singleton class with references to various other components (e.g. the `Catalog` and `BufferPool`). The tables are stored in an aptly named `Catalog` which keeps track of the relationship between the tables, their names, their ids, and their primary keys (if they are present). While interacting with SimpleDB, the `Database` is the appropriate method for interconnecting the different components, as they are not stand alone.

One component is the `BufferPool`, which keeps track of the pages that have recently been accessed (i.e. it functions as a cache of pages). This component cannot stand alone because it does not have its own method for reading the pages from disk. This is where the `HeapFile` comes in. It has a method for reading pages from disk into itself (since it is a representation of a file within the database). Once it reads a page, it can provide it to the `BufferPool` that can cache the page for quicker retrieval later.

The `HeapPage` is the random collection of `Tuples` that is used by the `BufferPool` and `HeapFile` as an intermediary to access the `Tuples`. Similarly, the `HeapPageId` just keeps track of which table and page number a specific `HeapPage` is tied to.

Finally, SimpleDB has operations which are handled by various types of operators. In this lab we worked through `SeqScan` which in most cases functions as a wrapper of the iterator returned by `HeapFile#iterator()`. This shows another interconnection between the components and demonstrates the higher level of abstraction that the clients would be more likely to interact with.

**Design decisions:**

- Tuple utilizes an underlying List<Field> for quick get/set of the data using indices.
- TupleDesc utilizes List<TDItem> for quick get/set of field names/types.
- Catalog utilizes two Maps with id→name and name→File/PrimaryKey. This is important because it allows the files to be accessed from both name and id. This also allows for quick lookup of the name from the id.
- BufferPool utilizes a Map with PageId→Page in conjunction with a maxCapacity value. This is important because it allows the BufferPool to cache a certain number of Pages that have been seen with quick lookup using their PageId.
- HeapFile utilizes a RandomAccessFile to deserialize the bytes from disk back into a HeapPage that can be passed to other components like the BufferPool. RandomAccessFile is important because it enables reading from an offset (i.e. the PageId that dictates the location of the serialized page).
- HeapFile#iterator() iterates Tuple by Tuple loading one Page in at a time (to prevent overloading the memory with larger tables). Additionally it stores the state of the iterator in the given HeapFile, so that when referenced by the Database singleton it's consistent.
- SeqScan#getTupleDesc() breaks the current TupleDesc stored for the given tableid into subparts (i.e. names and types) so that the alias can be prepended to each Field.

**Example of a unit test that could be added to improve the set:**

One unit test that could be added to improve the testing suite would be to verify that the BufferPool worked properly. For example, a unit test that asserts that the behavior of the permissions of the Page returned from BufferPool#getPage() is correct.

**Changes you made to the API:** None.

**Describe any missing or incomplete elements of your code:** None.

**Additional Feedback:** None.